


# DataCollider: Effective Data-Race Detection for the Kernel

刘伟森 PB15111595

王泽凡 PB15111593

- 
- ▶ Challenges for old methods
  - ▶ Insights
  - ▶ Implement
  - ▶ Advantage and Disadvantage
  - ▶ Result

# Challenges for old methods

- ▶ require a complete knowledge and logging of all locking semantics
- ▶ Locking semantics in kernel-mode can be complicated and convoluted
  - ▶ e.g. DPCs, interrupts

# Insights

- ▶ Instead of inferring if a data race could have occurred, let's cause it to actually happen!
- ▶ Use code and data breakpoints
- ▶ Randomly selection for uniform coverage

# Implement

- ▶ Sampling
- ▶ Insert code breakpoint
- ▶ Detection


# Detection

- ▶ Data breakpoint
- ▶ Repeated reads

# Detection

```
temp = read( loc, size );  
if ( isWrite )  
    SetDataBreakpointRW( loc, size );  
else  
    SetDataBreakpointW( loc, size );  
  
delay();  
  
ClearDataBreakpoint( loc, size );  
  
temp' = read( loc, size );  
if(temp != temp' || data breakpoint hit)  
    ReportDataRace( );
```


# Intersection Metaphor



@[OSDI 2010] DataCollider  
John Erickson



# Intersection Metaphor



# Intersection Metaphor


Instruction stream



Memory  
Address = 0x1000


@[OSDI 2010] DataCollider  
John Erickson

# Intersection Metaphor




@[OSDI 2010] DataCollider  
John Erickson


# Intersection Metaphor




# Intersection Metaphor



# Intersection Metaphor




# Intersection Metaphor





# Intersection Metaphor





# Intersection Metaphor: Normal Case

@[OSDI 2010] **DataCollider**  
John Erickson


# Intersection Metaphor: Normal Case



Thread B



DataCollider



Data Breakpoint

@[OSDI 2010] DataCollider  
John Erickson

# Intersection Metaphor: Normal Case



I don't have the lock,  
so I'll have to wait.




Data Breakpoint

DataCollider


@[OSDI 2010] DataCollider  
John Erickson

# Intersection Metaphor: Normal Case






# Intersection Metaphor: Normal Case



# Intersection Metaphor: Normal Case



@[OSDI 2010] DataCollider  
John Erickson

# Intersection Metaphor: Data Race

@[OSDI 2010] **DataCollider**  
John Erickson


# Intersection Metaphor: Data Race



Thread B



DataCollider




Data Breakpoint


@[OSDI 2010] DataCollider  
John Erickson



# Intersection Metaphor: Data Race




# Intersection Metaphor: Data Race



DataCollider


@[OSDI 2010] DataCollider  
John Erickson

# Intersection Metaphor: Data Race




@[OSDI 2010] DataCollider  
John Erickson

# Intersection Metaphor: Data Race



# Intersection Metaphor: Data Race



@[OSDI 2010] DataCollider  
John Erickson

# Data breakpoint

- ▶ Advantage
  - ▶ Setting the data breakpoint will catch the colliding thread in the act .
  - ▶ This provides much more actionable debugging information.
- ▶ Disadvantage
  - ▶ Works on virtual address

# Repeated reads

- ▶ Advantage

- ▶ The additional approach helps detect races caused by:
  - ▶ Hardware interaction via DMA
  - ▶ Physical memory that has multiple virtual mappings

- ▶ Disadvantage

- ▶ Cannot detect:
  - ▶ read conflicts at a breakpoint of write operation
  - ▶ Multi-writes but the value doesn't change

# Results: bugs found

- \* 25 confirmed bugs in the Windows OS have been found
- \* 8 more are still pending investigation

| Data Races Reported       | Count |
|---------------------------|-------|
| Fixed                     | 12    |
| Confirmed and Being Fixed | 13    |
| Under Investigation       | 8     |
| Harmless                  | 5     |
| Total                     | 38    |



# Some Problems

- ▶ Causing kernel crash?
- ▶ Sampling in other methods?
- ▶ Using DataCollider in user programs?

Thanks!